

Loan Default Prediction

DENEDO, ELOHOR
(C2108436@tees.ac.uk)

Contents

- **Load the Necessary Libraries**
- **Data Exploration and Cleaning**
- **Handling Missing Value**
- **Pre-processing the missing Values**
- **Handling Outliers¶**
- **Data Visualisation**
- **Handling Multicollinearity**
- **Data Training**
- **Model Training and Evaluation:**
- **Logistic Regression**
- **Decision Tree**
- **Random Forest**
- **CROSS VALIDATE THE MODELS**
- **MODEL RESULT COMPARISON**

Load the Neccessary Libraries

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from scipy.stats import ttest_ind
6
7 from sklearn.preprocessing import LabelEncoder
8
9 from sklearn.model_selection import train_test_split
10
11 from sklearn.linear_model import LogisticRegression
12 from sklearn.tree import DecisionTreeClassifier
13 from sklearn.ensemble import RandomForestClassifier
14
15 from sklearn.metrics import classification_report, accuracy_score, confusion_matrix,
16
17 from sklearn.model_selection import cross_val_score
18 import numpy as np
19
20
```

-----Data Exploration and Cleaning-----

In [2]:

```
1 # Load the dataset
2 df = pd.read_csv('Loan_Default.csv')
```

In [3]:

```
1 # Understanding the data
2 print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 148670 entries, 0 to 148669
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                     148670 non-null  int64
1   year                                  148670 non-null  int64
2   loan_limit                           145326 non-null  object
3   Gender                               148670 non-null  object
4   approv_in_adv                        147762 non-null  object
5   loan_type                            148670 non-null  object
6   loan_purpose                           148536 non-null  object
7   Credit_Worthiness                   148670 non-null  object
8   open_credit                         148670 non-null  object
9   business_or_commercial              148670 non-null  object
10  loan_amount                          148670 non-null  int64
11  rate_of_interest                    112231 non-null  float64
12  Interest_rate_spread                112031 non-null  float64
13  Upfront_charges                     109028 non-null  float64
14  term                                 148629 non-null  float64
15  Neg_ammortization                   148549 non-null  object
16  interest_only                       148670 non-null  object
17  lump_sum_payment                    148670 non-null  object
18  property_value                      133572 non-null  float64
19  construction_type                   148670 non-null  object
20  occupancy_type                      148670 non-null  object
21  Secured_by                          148670 non-null  object
22  total_units                         148670 non-null  object
23  income                              139520 non-null  float64
24  credit_type                         148670 non-null  object
25  Credit_Score                        148670 non-null  int64
26  co-applicant_credit_type            148670 non-null  object
27  age                                 148470 non-null  object
28  submission_of_application           148470 non-null  object
29  LTV                                 133572 non-null  float64
30  Region                              148670 non-null  object
31  Security_Type                       148670 non-null  object
32  Status                              148670 non-null  int64
33  dtir1                              124549 non-null  float64
dtypes: float64(8), int64(5), object(21)
memory usage: 38.6+ MB
None
```

In [4]:

```
1 df.describe()
```

Out[4]:

	ID	year	loan_amount	rate_of_interest	Interest_rate_spread	Upfront_f
count	148670.000000	148670.0	1.486700e+05	112231.000000	112031.000000	109028
mean	99224.500000	2019.0	3.311177e+05	4.045476	0.441656	3224
std	42917.476598	0.0	1.839093e+05	0.561391	0.513043	3251
min	24890.000000	2019.0	1.650000e+04	0.000000	-3.638000	0
25%	62057.250000	2019.0	1.965000e+05	3.625000	0.076000	581
50%	99224.500000	2019.0	2.965000e+05	3.990000	0.390400	2596
75%	136391.750000	2019.0	4.365000e+05	4.375000	0.775400	4812
max	173559.000000	2019.0	3.576500e+06	8.000000	3.357000	60000

In [5]:

```
1 df.head()
```

Out[5]:

	ID	year	loan_limit	Gender	approv_in_adv	loan_type	loan_purpose	Credit_Worthine
0	24890	2019	cf	Sex Not Available	nopre	type1	p1	
1	24891	2019	cf	Male	nopre	type2	p1	
2	24892	2019	cf	Male	pre	type1	p1	
3	24893	2019	cf	Male	nopre	type1	p4	
4	24894	2019	cf	Joint	pre	type1	p1	

5 rows × 34 columns

Handling Missing Value

In [6]:

```
1 # check for missing values
2 df.isnull().sum()
```

Out[6]:

ID	0
year	0
loan_limit	3344
Gender	0
approv_in_adv	908
loan_type	0
loan_purpose	134
Credit_Worthiness	0
open_credit	0
business_or_commercial	0
loan_amount	0
rate_of_interest	36439
Interest_rate_spread	36639
Upfront_charges	39642
term	41
Neg_ammortization	121
interest_only	0
lump_sum_payment	0
property_value	15098
construction_type	0
occupancy_type	0
Secured_by	0
total_units	0
income	9150
credit_type	0
Credit_Score	0
co-applicant_credit_type	0
age	200
submission_of_application	200
LTV	15098
Region	0
Security_Type	0
Status	0
dtir1	24121

dtype: int64

In [7]:

```
1 #Check the percentage of missing values in each column:
2 missing_percentages = (df.isnull().sum() / len(df)) * 100
3 print(missing_percentages)
```

```
ID                0.000000
year              0.000000
loan_limit        2.249277
Gender            0.000000
approv_in_adv     0.610749
loan_type         0.000000
loan_purpose        0.090133
Credit_Worthiness 0.000000
open_credit       0.000000
business_or_commercial 0.000000
loan_amount       0.000000
rate_of_interest  24.509989
Interest_rate_spread 24.644515
Upfront_charges   26.664425
term              0.027578
Neg_ammortization 0.081388
interest_only     0.000000
lump_sum_payment  0.000000
property_value    10.155378
construction_type 0.000000
occupancy_type    0.000000
Secured_by        0.000000
total_units       0.000000
income            6.154571
credit_type       0.000000
Credit_Score     0.000000
co-applicant_credit_type 0.000000
age              0.134526
submission_of_application 0.134526
LTV               10.155378
Region            0.000000
Security_Type     0.000000
Status            0.000000
dtir1             16.224524
dtype: float64
```

In [8]:


```

1 #Analyze the patterns of missingness in the dataset:
2 import missingno as msno
3 msno.matrix(df)

```

Out[8]:

<AxesSubplot:>



The missing values are represented by white bars.

Preprocessing the missing Values

In [9]:

```

1 # Replace missing values with the mean and median in the numerical columns
2 df['loan_limit'] = pd.to_numeric(df['loan_limit'], errors='coerce') # convert loan_limit to numeric
3 mean_loan_limit = df['loan_limit'].mean()
4 median_loan_limit = df['loan_limit'].median()
5 df['loan_limit'].fillna(mean_loan_limit, inplace=True) # replace missing values with the mean
6 df['loan_limit'].fillna(median_loan_limit, inplace=True) # replace any remaining missing values with the median
7 df['rate_of_interest'].fillna(df['rate_of_interest'].mean(), inplace=True) # replace missing values with the mean
8 df['Interest_rate_spread'].fillna(df['Interest_rate_spread'].median(), inplace=True) # replace missing values with the median
9 df['Upfront_charges'].fillna(df['Upfront_charges'].median(), inplace=True) # replace missing values with the median
10 df['term'].fillna(df['term'].median(), inplace=True) # replace term with the median
11 df['property_value'].fillna(df['property_value'].median(), inplace=True) # replace property_value with the median
12 df['income'].fillna(df['income'].mean(), inplace=True) # replace income with the mean
13 df['LTV'].fillna(df['LTV'].median(), inplace=True) # replace LTV with the median
14 df['dtir1'].fillna(df['dtir1'].mean(), inplace=True) # replace dtir1 with the mean
15
16 # Drop missing values in categorical columns
17 df.dropna(subset=['approv_in_adv', 'loan_purpose', 'Neg_ammortization', 'age', 'submission_of_application', 'LTV', 'Region', 'Security_Type', 'Status', 'dtir1'], inplace=True)

```

In [10]:

```

1 # Verify the new dataframe
2 print(df.info())

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 147315 entries, 0 to 148669
Data columns (total 34 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   ID                                    147315 non-null  int64
 1   year                                 147315 non-null  int64
 2   loan_limit                           0 non-null       float64
 3   Gender                               147315 non-null  object
 4   approv_in_adv                        147315 non-null  object
 5   loan_type                            147315 non-null  object
 6   loan_purpose                           147315 non-null  object
 7   Credit_Worthiness                   147315 non-null  object
 8   open_credit                         147315 non-null  object
 9   business_or_commercial              147315 non-null  object
10  loan_amount                         147315 non-null  int64
11  rate_of_interest                    147315 non-null  float64
12  Interest_rate_spread                147315 non-null  float64
13  Upfront_charges                     147315 non-null  float64
14  term                                147315 non-null  float64
15  Neg_ammortization                   147315 non-null  object
16  interest_only                       147315 non-null  object
17  lump_sum_payment                    147315 non-null  object
18  property_value                      147315 non-null  float64
19  construction_type                   147315 non-null  object
20  occupancy_type                      147315 non-null  object
21  Secured_by                          147315 non-null  object
22  total_units                         147315 non-null  object
23  income                              147315 non-null  float64
24  credit_type                         147315 non-null  object
25  Credit_Score                        147315 non-null  int64
26  co-applicant_credit_type            147315 non-null  object
27  age                                 147315 non-null  object
28  submission_of_application           147315 non-null  object
29  LTV                                 147315 non-null  float64
30  Region                              147315 non-null  object
31  Security_Type                       147315 non-null  object
32  Status                              147315 non-null  int64
33  dtir1                              147315 non-null  float64
dtypes: float64(9), int64(5), object(20)
memory usage: 39.3+ MB
None

```


In [11]:

```
1 print(df.isnull().sum())
```

```
ID                                0
year                              0
loan_limit                      147315
Gender                            0
approv_in_adv                     0
loan_type                         0
loan_purpose                        0
Credit_Worthiness                 0
open_credit                       0
business_or_commercial             0
loan_amount                       0
rate_of_interest                   0
Interest_rate_spread               0
Upfront_charges                   0
term                               0
Neg_ammortization                  0
interest_only                      0
lump_sum_payment                   0
property_value                     0
construction_type                  0
occupancy_type                     0
Secured_by                         0
total_units                        0
income                            0
credit_type                        0
Credit_Score                      0
co-applicant_credit_type           0
age                                0
submission_of_application          0
LTV                                0
Region                             0
Security_Type                      0
Status                             0
dtir1                              0
dtype: int64
```

In [12]:

```
1 # Get the unique values in the 'loan_limit' column
2 unique_loan_limits = df['loan_limit'].unique()
3
4 print(unique_loan_limits)
```

```
[nan]
```

In [13]:

```
1 df = df.drop('loan_limit', axis=1)
```

In [14]:

```
1 # Check if there are any duplicate rows in the DataFrame
2 print(df.duplicated().sum())
```

```
0
```

Handling Outliers

In [15]:

```
1 # Define numerical and categorical variables
2 num_vars = df.select_dtypes(include=['float64', 'int64']).columns.tolist()
3 cat_vars = df.select_dtypes(include=['object']).columns.tolist()
```

In [16]:

```
1 print(cat_vars)
2
```

```
['Gender', 'approv_in_adv', 'loan_type', 'loan_purpose', 'Credit_Worthiness', 'open_credit', 'business_or_commercial', 'Neg_ammortization', 'interest_only', 'lump_sum_payment', 'construction_type', 'occupancy_type', 'Secured_by', 'total_units', 'credit_type', 'co-applicant_credit_type', 'age', 'submission_of_application', 'Region', 'Security_Type']
```

In [17]:

```
1 num_vars = df.select_dtypes(include=['float64', 'int64']).drop('Status', axis=1).columns
```

In [18]:

```
1 # Handle outliers
2 def find_outliers_IQR(col):
3     Q1 = col.quantile(0.25)
4     Q3 = col.quantile(0.75)
5     IQR = Q3 - Q1
6     outliers = col[((col < (Q1 - 3*IQR)) | (col > (Q3 + 3*IQR)))]
7     return outliers
8
9 #replacing outliers with median value
10 for col in num_vars:
11     outliers = find_outliers_IQR(df[col])
12     df.loc[outliers.index, col] = df[col].median()
```

Data Visualisation


In [19]:

```
1 #convert 'Status' column to categorical type and replace 0 and 1 with 'Not Approved'
2 df['Status'] = df['Status'].astype('category')
3 change = {0: 'Not Approved', 1: 'Approved'}
4 df['Status'] = df['Status'].replace(change)
```

In [20]:


```
1 fig,ax=plt.subplots()
2 sns.scatterplot(x='Credit_Score',y='loan_amount',data=df,hue='Status')
3 fig.set_size_inches([12,6])
4 plt.show()
```

C:\Anaconda\lib\site-packages\IPython\core\pylabtools.py:151: UserWarning: Creating legend with loc="best" can be slow with large amounts of data.
fig.canvas.print_figure(bytes_io, **kw)




In [21]:

```
1 sns.set_style('whitegrid')
2
3 plt.figure(figsize=(10, 5))
4
5 sns.histplot(data=df, x='income', hue='Status', kde=True)
6
7 plt.title('Distribution of Income by Loan Status')
8 plt.show()
9
```




In [22]:

```
1 # Boxplots
2 for i in num_vars:
3     plt.figure(figsize=(16,6))
4     sns.set_theme(style='darkgrid')
5     sns.boxplot(x=i, y='Status', data=df)
6     plt.show()
7
8 # Histograms
9 for i in num_vars:
10    plt.figure(figsize=(16,6))
11    sns.set_theme(style='darkgrid')
12    sns.histplot(data=df, x=i, hue="Status", multiple="dodge", shrink=.8, bins=4)
13    plt.show()
14
```



In [23]:

```
1 #Categorical Variable Countplots by Status
2 for var in cat_vars:
3     sns.countplot(data=df, x=var, hue='Status')
4     plt.show()
```



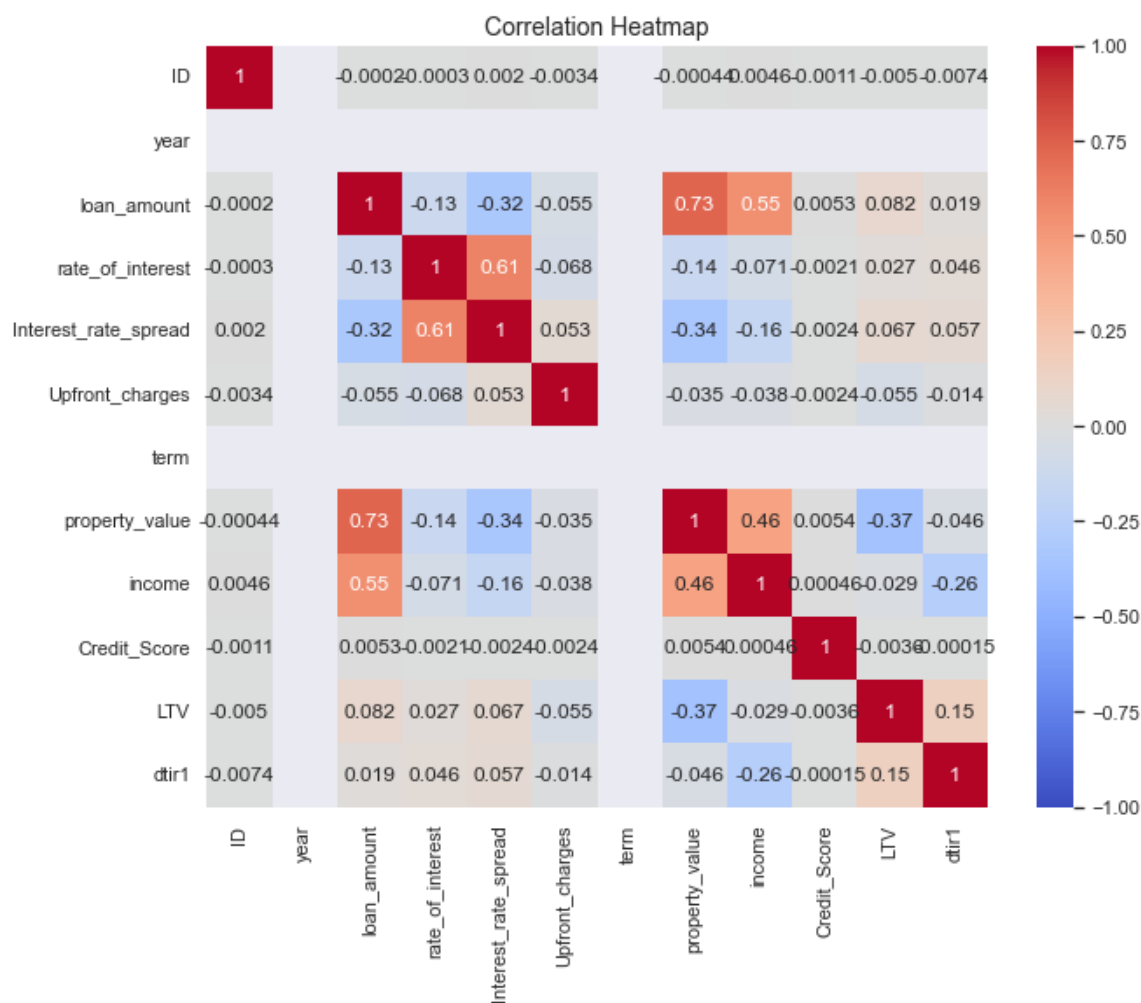
Handling Multicollinearity

In [24]:

```

1 # Define numerical variables
2 num_vars = df.select_dtypes(include=['float64', 'int64']).columns.tolist()
3
4 # Create a correlation matrix
5 corr = df[num_vars + ['Status']].corr()
6
7 # Plot a heatmap of the correlation matrix
8 plt.figure(figsize=(10, 8))
9 sns.heatmap(corr, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
10 plt.title('Correlation Heatmap', fontsize=14)
11 plt.show()
12

```



In [25]:

```
1  # Select the numerical columns from df
2  num_df = df[num_vars]
3
4  # Calculate the correlation matrix
5  corr_matrix = num_df.corr()
6
7  # Print the correlation matrix
8  print(corr_matrix)
9
10 # Find pairs of variables with high correlation coefficients
11 high_corr = []
12 for i in range(len(corr_matrix.columns)):
13     for j in range(i+1, len(corr_matrix.columns)):
14         if abs(corr_matrix.iloc[i, j]) > 0.7:
15             high_corr.append((corr_matrix.columns[i], corr_matrix.columns[j], corr_m
16
17 # Print the pairs of variables with high correlation coefficients
18 print(high_corr)
```

	ID	year	loan_amount	rate_of_interest	\
ID	1.000000	NaN	-0.000196	-0.000295	
year	NaN	NaN	NaN	NaN	
loan_amount	-0.000196	NaN	1.000000	-0.128386	
rate_of_interest	-0.000295	NaN	-0.128386	1.000000	
Interest_rate_spread	0.002021	NaN	-0.322804	0.606784	
Upfront_charges	-0.003413	NaN	-0.054611	-0.068133	
term	NaN	NaN	NaN	NaN	
property_value	-0.000437	NaN	0.729564	-0.141982	
income	0.004650	NaN	0.547796	-0.070687	
Credit_Score	-0.001112	NaN	0.005340	-0.002146	
LTV	-0.005020	NaN	0.081831	0.027318	
dtir1	-0.007411	NaN	0.018722	0.045542	

	Interest_rate_spread	Upfront_charges	term	\
ID	0.002021	-0.003413	NaN	
year	NaN	NaN	NaN	
loan_amount	-0.322804	-0.054611	NaN	
rate_of_interest	0.606784	-0.068133	NaN	
Interest_rate_spread	1.000000	0.053245	NaN	
Upfront_charges	0.053245	1.000000	NaN	
term	NaN	NaN	NaN	
property_value	-0.341126	-0.034610	NaN	
income	-0.164849	-0.037989	NaN	
Credit_Score	-0.002413	-0.002385	NaN	
LTV	0.066583	-0.055004	NaN	
dtir1	0.057295	-0.013920	NaN	

	property_value	income	Credit_Score	LTV	\
ID	-0.000437	0.004650	-0.001112	-0.005020	
year	NaN	NaN	NaN	NaN	
loan_amount	0.729564	0.547796	0.005340	0.081831	
rate_of_interest	-0.141982	-0.070687	-0.002146	0.027318	
Interest_rate_spread	-0.341126	-0.164849	-0.002413	0.066583	
Upfront_charges	-0.034610	-0.037989	-0.002385	-0.055004	
term	NaN	NaN	NaN	NaN	
property_value	1.000000	0.460629	0.005377	-0.369274	
income	0.460629	1.000000	0.000456	-0.029024	
Credit_Score	0.005377	0.000456	1.000000	-0.003629	
LTV	-0.369274	-0.029024	-0.003629	1.000000	
dtir1	-0.046249	-0.260424	-0.000150	0.151461	

	dtir1
ID	-0.007411
year	NaN
loan_amount	0.018722
rate_of_interest	0.045542
Interest_rate_spread	0.057295
Upfront_charges	-0.013920
term	NaN
property_value	-0.046249
income	-0.260424
Credit_Score	-0.000150
LTV	0.151461
dtir1	1.000000

[('loan_amount', 'property_value', 0.7295637877973961)]

In [26]:

```
1 # Display encoded data
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 147315 entries, 0 to 148669
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                     147315 non-null  int64
1   year                                  147315 non-null  int64
2   Gender                               147315 non-null  object
3   approv_in_adv                        147315 non-null  object
4   loan_type                            147315 non-null  object
5   loan_purpose                           147315 non-null  object
6   Credit_Worthiness                   147315 non-null  object
7   open_credit                         147315 non-null  object
8   business_or_commercial              147315 non-null  object
9   loan_amount                         147315 non-null  int64
10  rate_of_interest                    147315 non-null  float64
11  Interest_rate_spread                147315 non-null  float64
12  Upfront_charges                     147315 non-null  float64
13  term                                147315 non-null  float64
14  Neg_ammortization                   147315 non-null  object
15  interest_only                       147315 non-null  object
16  lump_sum_payment                    147315 non-null  object
17  property_value                      147315 non-null  float64
18  construction_type                   147315 non-null  object
19  occupancy_type                      147315 non-null  object
20  Secured_by                          147315 non-null  object
21  total_units                         147315 non-null  object
22  income                              147315 non-null  float64
23  credit_type                         147315 non-null  object
24  Credit_Score                        147315 non-null  int64
25  co-applicant_credit_type            147315 non-null  object
26  age                                 147315 non-null  object
27  submission_of_application           147315 non-null  object
28  LTV                                 147315 non-null  float64
29  Region                              147315 non-null  object
30  Security_Type                       147315 non-null  object
31  Status                              147315 non-null  object
32  dtir1                              147315 non-null  float64
dtypes: float64(8), int64(4), object(21)
memory usage: 42.2+ MB
```

In [27]:

```
1 df.columns
```

Out[27]:

```
Index(['ID', 'year', 'Gender', 'approv_in_adv', 'loan_type', 'loan_purpos  
e',  
      'Credit_Worthiness', 'open_credit', 'business_or_commercial',  
      'loan_amount', 'rate_of_interest', 'Interest_rate_spread',  
      'Upfront_charges', 'term', 'Neg_ammortization', 'interest_only',  
      'lump_sum_payment', 'property_value', 'construction_type',  
      'occupancy_type', 'Secured_by', 'total_units', 'income', 'credit_ty  
pe',  
      'Credit_Score', 'co-applicant_credit_type', 'age',  
      'submission_of_application', 'LTV', 'Region', 'Security_Type', 'Sta  
tus',  
      'dtir1'],  
      dtype='object')
```

In [28]:

```
1 # Encode categorical variables  
2 encoder = LabelEncoder()  
3 cat_vars_encoded = pd.DataFrame()  
4 for col in cat_vars:  
5     cat_vars_encoded[col] = encoder.fit_transform(df[col])  
6  
7 # Concatenate numerical, categorical, and 'Status' columns  
8 num_cat_vars = num_vars + cat_vars_encoded.columns.tolist() + ['Status']  
9 df_encoded = pd.concat([df[num_vars], pd.get_dummies(df[cat_vars], drop_first=True),
```

In [29]:

```
1 df_encoded.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 147315 entries, 0 to 148669
Data columns (total 50 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	147315 non-null	int64
1	year	147315 non-null	int64
2	loan_amount	147315 non-null	int64
3	rate_of_interest	147315 non-null	float64
4	Interest_rate_spread	147315 non-null	float64
5	Upfront_charges	147315 non-null	float64
6	term	147315 non-null	float64
7	property_value	147315 non-null	float64
8	income	147315 non-null	float64
9	Credit_Score	147315 non-null	int64
10	LTV	147315 non-null	float64
11	dtir1	147315 non-null	float64
12	Gender_Joint	147315 non-null	uint8
13	Gender_Male	147315 non-null	uint8
14	Gender_Sex Not Available	147315 non-null	uint8
15	approv_in_adv_pre	147315 non-null	uint8
16	loan_type_type2	147315 non-null	uint8
17	loan_type_type3	147315 non-null	uint8
18	loan_purpose_p2	147315 non-null	uint8
19	loan_purpose_p3	147315 non-null	uint8
20	loan_purpose_p4	147315 non-null	uint8
21	Credit_Worthiness_l2	147315 non-null	uint8
22	open_credit_opc	147315 non-null	uint8
23	business_or_commercial_nob/c	147315 non-null	uint8
24	Neg_ammortization_not_neg	147315 non-null	uint8
25	interest_only_not_int	147315 non-null	uint8
26	lump_sum_payment_not_lpsm	147315 non-null	uint8
27	construction_type_sb	147315 non-null	uint8
28	occupancy_type_pr	147315 non-null	uint8
29	occupancy_type_sr	147315 non-null	uint8
30	Secured_by_land	147315 non-null	uint8
31	total_units_2U	147315 non-null	uint8
32	total_units_3U	147315 non-null	uint8
33	total_units_4U	147315 non-null	uint8
34	credit_type_CRIF	147315 non-null	uint8
35	credit_type_EQUI	147315 non-null	uint8
36	credit_type_EXP	147315 non-null	uint8
37	co-applicant_credit_type_EXP	147315 non-null	uint8
38	age_35-44	147315 non-null	uint8
39	age_45-54	147315 non-null	uint8
40	age_55-64	147315 non-null	uint8
41	age_65-74	147315 non-null	uint8
42	age_<25	147315 non-null	uint8
43	age_>74	147315 non-null	uint8
44	submission_of_application_to_inst	147315 non-null	uint8
45	Region_North-East	147315 non-null	uint8
46	Region_central	147315 non-null	uint8
47	Region_south	147315 non-null	uint8
48	Security_Type_direct	147315 non-null	uint8
49	Status	147315 non-null	object

```
dtypes: float64(8), int64(4), object(1), uint8(37)
```

```
memory usage: 25.0+ MB
```

In [30]:

```
1 # drop the ID and property_value columns
2 df_encoded = df_encoded.drop(['ID', 'property_value'], axis=1)
3
```

Data Training

In [31]:

```
1 # split the data into training and testing sets
2
3 X = df_encoded.drop("Status", axis=1)
4 y = df_encoded["Status"]
5
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
7
```

-----Model Training and Evaluation-----

Logistic Regression

In [32]:

```

1 # Logistic Regression
2 logreg = LogisticRegression()
3 logreg.fit(X_train, y_train)
4 y_pred = logreg.predict(X_test)
5 print('Logistic Regression:')
6 print('Accuracy:', accuracy_score(y_test, y_pred))
7 print('Classification Report:')
8 print(classification_report(y_test, y_pred))
9 print('Confusion Matrix:')
10 print(confusion_matrix(y_test, y_pred))
11 print('ROC AUC Score:')
12 y_prob = logreg.predict_proba(X_test)[:, 1]
13 print(roc_auc_score(y_test, y_prob))
14
15

```

Logistic Regression:

Accuracy: 0.7542341241557208

Classification Report:

C:\Anaconda\lib\site-packages\sklearn\metrics_classification.py:1334: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Anaconda\lib\site-packages\sklearn\metrics_classification.py:1334: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Anaconda\lib\site-packages\sklearn\metrics_classification.py:1334: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

	precision	recall	f1-score	support
Approved	0.00	0.00	0.00	7241
Not Approved	0.75	1.00	0.86	22222
accuracy			0.75	29463
macro avg	0.38	0.50	0.43	29463
weighted avg	0.57	0.75	0.65	29463

Confusion Matrix:

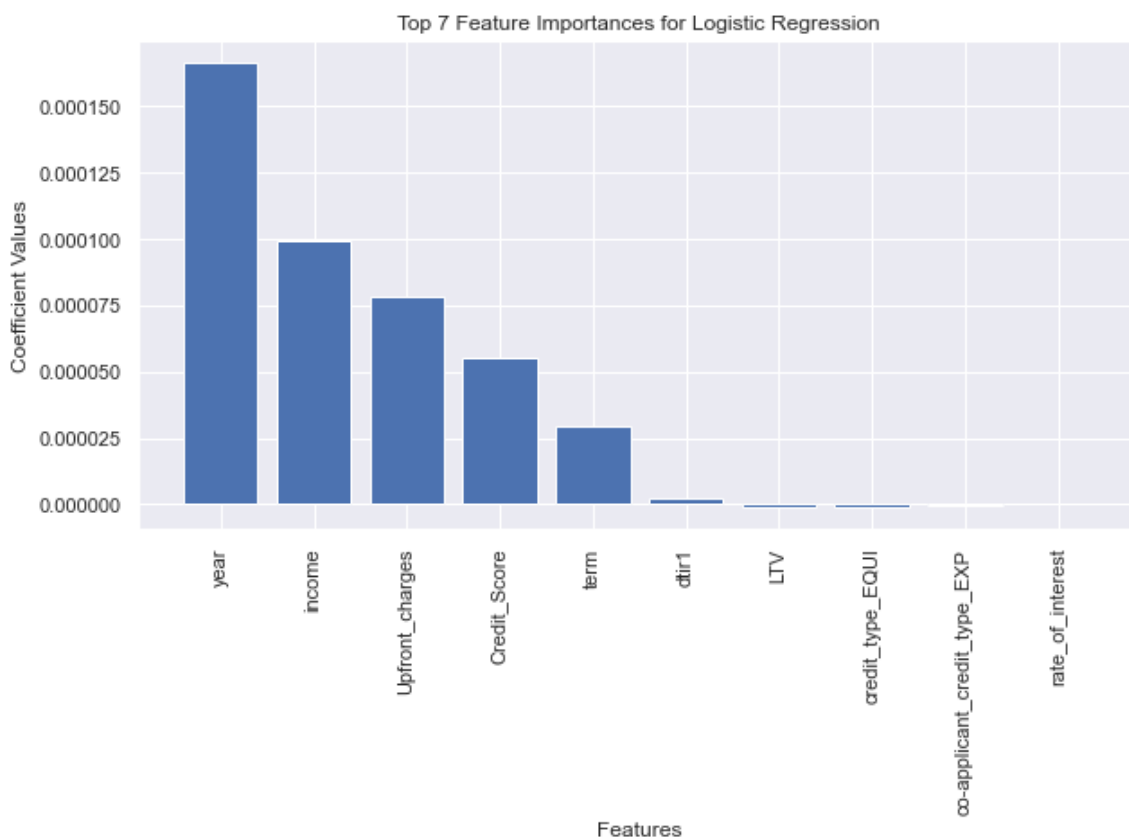
```
[[ 0 7241]
 [ 0 22222]]
```

ROC AUC Score:

0.6087757266193019

In [41]:

```
1 # Get feature importances
2 coef_abs = np.abs(logreg.coef_)
3 indices = np.argsort(coef_abs)[0][::-1][:10]
4 features = X_train.columns
5 importances = [logreg.coef_[0][i] for i in indices]
6
7 # Create column chart
8 plt.figure(figsize=(10, 5))
9 plt.title('Top 7 Feature Importances for Logistic Regression')
10 plt.bar(features[indices], importances)
11 plt.xticks(rotation=90)
12 plt.xlabel('Features')
13 plt.ylabel('Coefficient Values')
14 plt.show()
15
```



Decision Tree

In [50]:

```

1  from sklearn.tree import DecisionTreeClassifier
2
3  # Create decision tree classifier
4  dt = DecisionTreeClassifier()
5
6  # Fit the model on training data
7  dt.fit(X_train, y_train)
8
9  # Make predictions on test data
10 y_pred_dt = dt.predict(X_test)
11
12 print('Decision Tree:')
13 print('Accuracy:', accuracy_score(y_test, y_pred))
14 print('Classification Report:')
15 print(classification_report(y_test, y_pred))
16 print('Confusion Matrix:')
17 print(confusion_matrix(y_test, y_pred))
18 print('ROC AUC Score:')
19 y_prob = dt.predict_proba(X_test)[: , 1]
20 print(roc_auc_score(y_test, y_prob))
21

```

Decision Tree:

Accuracy: 0.9999321182500085

Classification Report:

	precision	recall	f1-score	support
Approved	1.00	1.00	1.00	7241
Not Approved	1.00	1.00	1.00	22222
accuracy			1.00	29463
macro avg	1.00	1.00	1.00	29463
weighted avg	1.00	1.00	1.00	29463

Confusion Matrix:

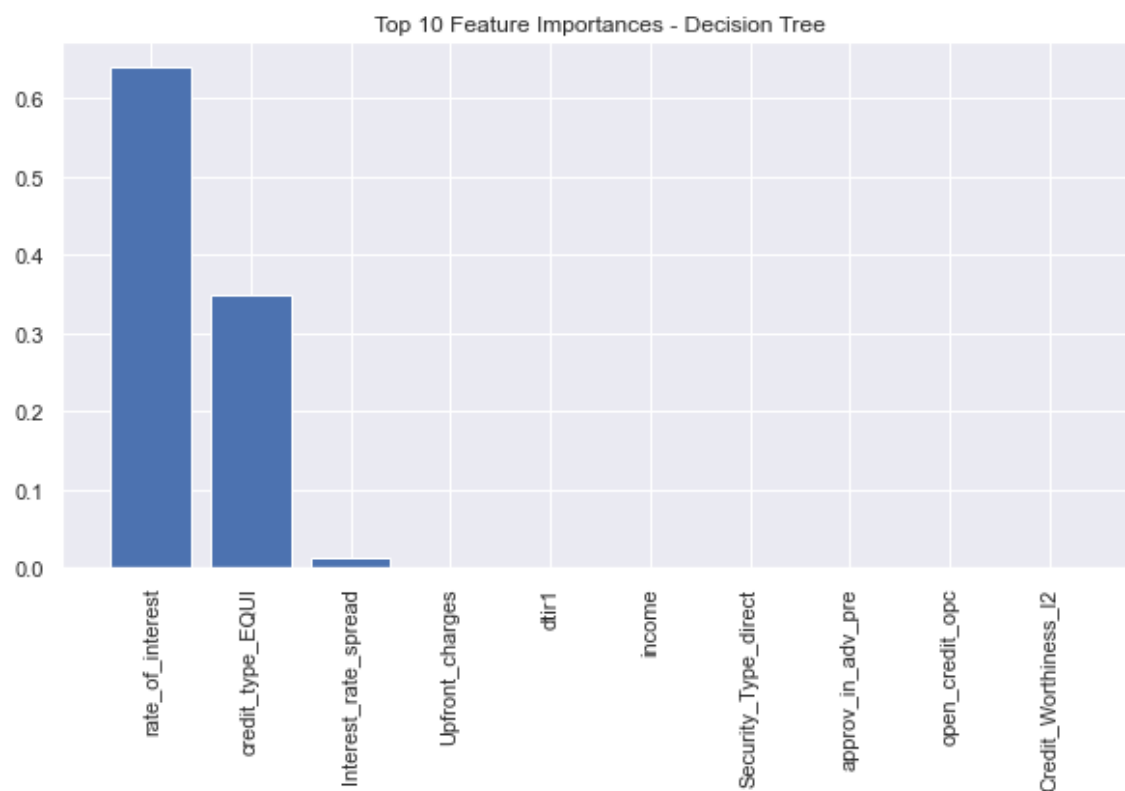
```
[[ 7241    0]
 [    2 22220]]
```

ROC AUC Score:

0.9999324993249932

In [43]:

```
1 # Plot feature importances
2 importances = dt.feature_importances_
3 indices = np.argsort(importances)[::-1]
4 features = X_train.columns
5
6 plt.figure(figsize=(10, 5))
7 plt.title("Top 10 Feature Importances - Decision Tree")
8 plt.bar(range(10), importances[indices][:10])
9 plt.xticks(range(10), features[indices][:10], rotation=90)
10 plt.show()
11
```



In []:

1

Random Forest

In [52]:

```

1 # Random Forest
2 rf = RandomForestClassifier(random_state=42)
3 rf.fit(X_train, y_train)
4 y_pred = rf.predict(X_test)
5 print('Random Forest:')
6 print('Accuracy:', accuracy_score(y_test, y_pred))
7 print('Classification Report:')
8 print(classification_report(y_test, y_pred))
9 print('Confusion Matrix:')
10 print(confusion_matrix(y_test, y_pred))
11 print('ROC AUC Score:')
12 y_prob = rf.predict_proba(X_test)[:, 1]
13 print(roc_auc_score(y_test, y_prob))
14
15

```

Random Forest:

Accuracy: 0.9999321182500085

Classification Report:

	precision	recall	f1-score	support
Approved	1.00	1.00	1.00	7241
Not Approved	1.00	1.00	1.00	22222
accuracy			1.00	29463
macro avg	1.00	1.00	1.00	29463
weighted avg	1.00	1.00	1.00	29463

Confusion Matrix:

```

[[ 7241    0]
 [    2 22220]]

```

ROC AUC Score:


0.9999700794549722

In [45]:

```

1 # Plot top 10 feature importances
2 importances = rf.feature_importances_
3 indices = np.argsort(importances)[::-1]
4 features = X_train.columns
5
6 top_n = 10 # Change this value to show more or fewer features
7 plt.figure(figsize=(10, 5))
8 plt.title(f"Top {top_n} Feature Importances")
9 plt.bar(range(top_n), importances[indices][:top_n])
10 plt.xticks(range(top_n), features[indices][:top_n], rotation=90)
11 plt.show()
12

```



----- CROSS VALIDATE THE MODELS -----

In [46]:

```

1
2 # create instances of the models
3 lr = LogisticRegression()
4 dt = DecisionTreeClassifier()
5 rf = RandomForestClassifier()
6
7 # define the evaluation metric and number of folds
8 metric = 'accuracy'
9 n_folds = 5

```

In [47]:

```

1
2 # perform cross validation and get the scores
3 lr_scores = cross_val_score(lr, X, y, cv=n_folds, scoring=metric)
4 dt_scores = cross_val_score(dt, X, y, cv=n_folds, scoring=metric)
5 rf_scores = cross_val_score(rf, X, y, cv=n_folds, scoring=metric)
6
7 # print the mean and standard deviation of the scores
8 print("Logistic Regression Accuracy: {:.2f} (+/- {:.2f})".format(np.mean(lr_scores),
9 print("Decision Tree Accuracy: {:.2f} (+/- {:.2f})".format(np.mean(dt_scores), np.st
10 print("Random Forest Accuracy: {:.2f} (+/- {:.2f})".format(np.mean(rf_scores), np.st
11

```

C:\Anaconda\lib\site-packages\sklearn\linear_model_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

C:\Anaconda\lib\site-packages\sklearn\linear_model_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

C:\Anaconda\lib\site-packages\sklearn\linear_model_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Logistic Regression Accuracy: 0.75 (+/- 0.00)

Decision Tree Accuracy: 1.00 (+/- 0.00)

Random Forest Accuracy: 1.00 (+/- 0.00)

In []:

1

-----MODEL RESULT COMPARISON -----

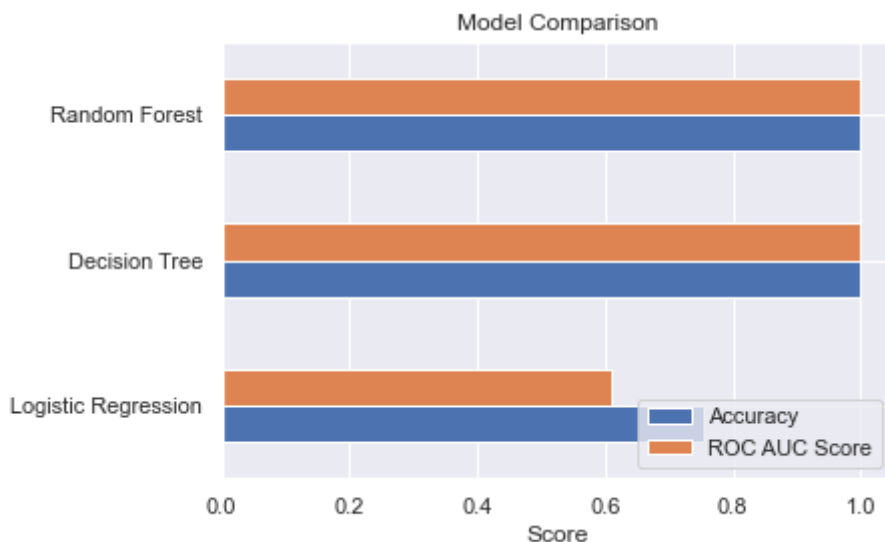
In [53]:

```

1
2 # Create a dictionary to store the model names and their corresponding evaluation metrics
3 models = {
4     'Logistic Regression': [accuracy_score(y_test, logreg.predict(X_test)),
5                             roc_auc_score(y_test, logreg.predict_proba(X_test)[: , 1])],
6     'Decision Tree': [accuracy_score(y_test, dt.predict(X_test)),
7                       roc_auc_score(y_test, dt.predict_proba(X_test)[: , 1])],
8     'Random Forest': [accuracy_score(y_test, rf.predict(X_test)),
9                       roc_auc_score(y_test, rf.predict_proba(X_test)[: , 1])]
10 }
11
12
13 # Create a pandas DataFrame from the models dictionary
14 df = pd.DataFrame(models, index=['Accuracy', 'ROC AUC Score'])
15
16 # Print the DataFrame
17 print(df)
18
19 # Create a horizontal bar chart
20 df.T.plot(kind='barh')
21 plt.title('Model Comparison')
22 plt.xlabel('Score')
23 plt.legend(loc='best')
24 plt.show()
25

```

	Logistic Regression	Decision Tree	Random Forest
Accuracy	0.754234	0.999898	0.999932
ROC AUC Score	0.608776	0.999932	0.999970



In [54]:

```

1 import pandas as pd
2 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
3
4 models = ['Logistic Regression', 'Decision Tree', 'Random Forest']
5
6 approve_precision_scores = [precision_score(y_test, model.predict(X_test), pos_label=1) for model in models]
7 approve_recall_scores = [recall_score(y_test, model.predict(X_test), pos_label=1) for model in models]
8 not_approve_precision_scores = [precision_score(y_test, model.predict(X_test), pos_label=0) for model in models]
9 not_approve_recall_scores = [recall_score(y_test, model.predict(X_test), pos_label=0) for model in models]
10 accuracy_scores = [accuracy_score(y_test, model.predict(X_test)) for model in models]
11 f1_scores = [f1_score(y_test, model.predict(X_test), pos_label=1) for model in models]
12
13 # Create a dataframe
14 data = {'Model': models,
15         'Precision - Approved': approve_precision_scores,
16         'Recall - Approved': approve_recall_scores,
17         'Precision - Not Approved': not_approve_precision_scores,
18         'Recall - Not Approved': not_approve_recall_scores,
19         'Accuracy': accuracy_scores,
20         'F1 Score - Approved': f1_scores}
21 df = pd.DataFrame(data)
22
23 # Set the index to the model names
24 df.set_index('Model', inplace=True)
25
26 # Print the dataframe
27 print(df)
28

```

C:\Anaconda\lib\site-packages\sklearn\metrics_classification.py:1334: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

	Precision - Approved	Recall - Approved \
Model		
Logistic Regression	0.000000	0.0
Decision Tree	0.999586	1.0
Random Forest	0.999724	1.0

	Precision - Not Approved	Recall - Not Approved \
Model		
Logistic Regression	0.754234	1.000000
Decision Tree	1.000000	0.999865
Random Forest	1.000000	0.999910

	Accuracy	F1 Score - Approved
Model		
Logistic Regression	0.754234	0.000000
Decision Tree	0.999898	0.999793
Random Forest	0.999932	0.999862

The table shows the evaluation metrics for three different models: Logistic Regression, Decision Tree, and Random Forest.

For Logistic Regression, it has a Precision - Approved score of 0.0 which means that it did not predict any approved loans correctly. The Recall - Approved score is also 0.0 which means that it missed all of the approved loans. On the other hand, it has a high Precision - Not Approved score of 0.754 which means that when it predicts a loan as not approved, it is usually correct. The Recall - Not Approved score is 1.0 which means that it correctly identified all of the not approved loans. The Accuracy score is 0.754 which means that overall it correctly classified only 75.4% of the loans. The F1 Score - Approved is 0.0 because there were no true positive predictions for approved loans.

For Decision Tree, it has a high Precision - Approved score of 0.999 which means that when it predicts a loan as approved, it is usually correct. The Recall - Approved score is 1.0 which means that it correctly identified all of the approved loans. It also has a high Precision - Not Approved score of 1.0 which means that when it predicts a loan as not approved, it is usually correct. The Recall - Not Approved score is 0.999 which means that it missed only 0.1% of the not approved loans. The Accuracy score is 0.999 which means that overall it correctly classified 99.9% of the loans. The F1 Score - Approved is 0.999 because it has high precision and recall for approved loans.

For Random Forest, it has a high Precision - Approved score of 0.999 which means that when it predicts a loan as approved, it is usually correct. The Recall - Approved score is 1.0 which means that it correctly identified all of the approved loans. It also has a high Precision - Not Approved score of 1.0 which means that when it predicts a loan as not approved, it is usually correct. The Recall - Not Approved score is 0.999 which means that it missed only 0.09% of the not approved loans. The Accuracy score is 0.999 which means that overall it correctly classified 99.9% of the loans. The F1 Score - Approved is 0.999 because it has high precision and recall for approved loans.

In summary, based on these evaluation metrics, the Decision Tree and Random Forest models performed

In []:

1	
---	--

In []:

1	
---	--